

## Exercise 7 : Bayes Optimal Classifier, Naïve Bayes, Random Variables and Distributions, EM Clustering

### Exercise 7-1 : Bayes Optimal

We have a classification problem with two classes “+” and “−”, three trained classifiers  $h_1$ ,  $h_2$ , and  $h_3$ , with the following probabilities of the classifiers, given the training data  $D$  :

$$\Pr(h_1|D) = 0.5$$

$$\Pr(h_2|D) = 0.3$$

$$\Pr(h_3|D) = 0.2$$

For the three test instances  $o_1$ ,  $o_2$ ,  $o_3$ , the classifiers give the following class probabilities :

$o_1 : \Pr(+ h_1) = 0.6$	$\Pr(- h_1) = 0.4$
$\Pr(+ h_2) = 0.2$	$\Pr(- h_2) = 0.8$
$\Pr(+ h_3) = 0.9$	$\Pr(- h_3) = 0.1$
$o_2 : \Pr(+ h_1) = 0.6$	$\Pr(- h_1) = 0.4$
$\Pr(+ h_2) = 0.6$	$\Pr(- h_2) = 0.4$
$\Pr(+ h_3) = 1$	$\Pr(- h_3) = 0$
$o_3 : \Pr(+ h_1) = 0.6$	$\Pr(- h_1) = 0.4$
$\Pr(+ h_2) = 0.6$	$\Pr(- h_2) = 0.4$
$\Pr(+ h_3) = 0$	$\Pr(- h_3) = 1$

We combine the three classifiers to get a Bayes optimal classifier. Which class probabilities will we get from this Bayes optimal classifier for the three test instances ?

**Exercise 7-2 : Naïve Bayes**

The skiing season is open. To reliably decide when to go skiing and when not, you could use a classifier such as Naïve Bayes. The classifier will be trained with your observations from the last year. Your notes include the following attributes :

The weather : The attribute **weather** can have the following three values : **sunny**, **rainy**, and **snow**.

The snow level : The attribute **snow level** can have the following two values :  $\geq 50$  (There are at least 50 cm of snow) and  $< 50$  (There are less than 50 cm of snow).

Assume you went skiing 8 times during the previous year. Here is the table with your decisions :

<b>weather</b>	<b>snow level</b>	<b>ski ?</b>
sunny	$< 50$	no
rainy	$< 50$	no
rainy	$\geq 50$	no
snow	$\geq 50$	yes
snow	$< 50$	no
sunny	$\geq 50$	yes
snow	$\geq 50$	yes
rainy	$< 50$	yes

- Compute the *a priori* probabilities for both classes **ski = yes** and **ski = no** (on the training set)!
- Compute the distribution of the conditional probabilities for the two classes for each attribute.
- Decide for the following weather and snow conditions, whether to go skiing or not! Use the Naïve Bayes classifier as trained in the previous steps for your decision.

	<b>weather</b>	<b>snow level</b>
day A	sunny	$\geq 50$
day B	rainy	$< 50$
day C	snow	$< 50$

**Exercise 7-3 : Assignments in the EM-Algorithm**

Given a data set with 100 points consisting of three Gaussian clusters  $A$ ,  $B$  and  $C$  and the point  $p$ .

The cluster  $A$  contains 30% of all objects and is represented using the mean of all its points  $\mu_A = (2, 2)$

and the covariance matrix  $\Sigma_A = \begin{pmatrix} 3 & 0 \\ 0 & 3 \end{pmatrix}$ .

You will need the inverse :  $\Sigma_A^{-1} = \begin{pmatrix} \frac{1}{3} & 0 \\ 0 & \frac{1}{3} \end{pmatrix}$ .

The cluster  $B$  contains 20% of all objects and is represented using the mean of all its points  $\mu_B = (5, 3)$


and the covariance matrix  $\Sigma_B = \begin{pmatrix} 2 & 1 \\ 1 & 4 \end{pmatrix}$ .  $\Sigma_B^{-1} \approx \begin{pmatrix} 0.571428 & -0.142857 \\ -0.142857 & 0.285714 \end{pmatrix}$ .

The cluster  $C$  contains 50% of all objects and is represented using the mean of all its points  $\mu_C = (1, 4)$

and the covariance matrix  $\Sigma_C = \begin{pmatrix} 16 & 0 \\ 0 & 4 \end{pmatrix}$ .  $\Sigma_C^{-1} = \begin{pmatrix} \frac{1}{16} & 0 \\ 0 & \frac{1}{4} \end{pmatrix}$ .

The point  $p$  is given by the coordinates  $(2.5, 3.0)$ .

The following sketch is not exact, and only gives a rough idea of the cluster locations :



Compute the three probabilities of  $p$  belonging to the clusters  $A$ ,  $B$ , and  $C$ .

**Exercise 7-4 : Tools : EM algorithm**

Consider EM algorithm on iris dataset as bellow :

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
import scipy.stats
import seaborn as sns
import numpy as np

# import some data to play with
iris = datasets.load_iris()
X = iris.data # we only take the first two features.
y = iris.target

X = PCA(n_components=2).fit_transform(iris.data)
# -----
# The EM algo
# -----

def normal_density(X, mu, Sigma):
    L = np.linalg.cholesky(Sigma)
    Linv = np.linalg.inv(L)
    Sinv = Linv.T.dot(Linv)

    XL = X.dot(Linv)
    # P stands for precision, i.e. inverse Sigma
    xPx = (XL*XL).sum(axis=1)
    xPmu = X.dot(Sinv).dot(mu)
    muPmu = mu.dot(Sinv).dot(mu)
    mahalanobis = xPx - 2*xPmu + muPmu
    twoPiPowD = (2*np.pi)**D
    sqrtDetSigma = L.diagonal().prod()
    density = 1/(np.sqrt(twoPiPowD)*sqrtDetSigma)*
               np.exp(-0.5*(mahalanobis))

    return density

K = 2 # Cluster count
max_iter = 20
(N,D) = X.shape

# Initialize
mu = np.random.randn(K,D)
```

```

Sigma = np.zeros([K,D,D])
for k in range(K):
    #L = np.random.randn(D,D)
    Sigma[k,:,:] = np.eye(D) #+ L.dot(L.T)
cls_prob = np.zeros([N,K])
pi_k = np.ones(K)/K

list_log_lik = np.zeros([max_iter])

for iter in range(max_iter):

    # E-STEP -----
    # Update cluster probabilities
    for k in range(K):
        cls_prob[:,k] = pi_k[k]*normal_density(X,mu[k,:],Sigma[k,:,:])
    cls_prob = cls_prob / np.broadcast_to(np.expand_dims
        (cls_prob.sum(axis=1), axis=1),(N,K))

    Nk = cls_prob.sum(axis=0)
    pi_k = Nk / Nk.sum()

    # M-STEP -----
    # Update means and covariances
    for k in range(K):
        clsProbMat = np.broadcast_to(np.expand_dims(cls_prob[:,k]
            ,axis=1),(N,D))

        mu[k,:] = 1/Nk[k]*(X*clsProbMat).sum(axis=0)

        Z = (X - mu[k,:])*np.sqrt(clsProbMat)
        ZtZ = 1/Nk[k]*Z.T.dot(Z)
        Sigma[k,:,:] = 1/Nk[k]*ZtZ + np.eye(D)

    # Report model fit -----
    evidence = 0
    for k in range(K):
        evidence += pi_k[k]*normal_density(X,mu[k,:],Sigma[k,:,:])
    list_log_lik[iter] = np.log(evidence).sum()

```

- Rerun the algorithm for different number of clusters.
- Plot marginal log likelihood and cluster assignments for different values of K.
- Plot heatmaps for different values of K.
- Describe which k can better fit the model.